

Lab 3 - Linear Regression

An Introduction to Statistical Learning

```
library(MASS)
library(ISLR)
attach(Boston)
```

```
summary(Boston)
```

```
##      crim      zn      indus      chas
## Min.   : 0.00632 Min.   : 0.00 Min.   : 0.46 Min.   :0.00000
## 1st Qu.: 0.08205 1st Qu.: 0.00 1st Qu.: 5.19 1st Qu.:0.00000
## Median : 0.25651 Median : 0.00 Median : 9.69 Median :0.00000
## Mean   : 3.61352 Mean   : 11.36 Mean   :11.14 Mean   :0.06917
## 3rd Qu.: 3.67708 3rd Qu.: 12.50 3rd Qu.:18.10 3rd Qu.:0.00000
## Max.   :88.97620 Max.   :100.00 Max.   :27.74 Max.   :1.00000
##      nox      rm      age      dis
## Min.   :0.3850 Min.   :3.561 Min.   : 2.90 Min.   : 1.130
## 1st Qu.:0.4490 1st Qu.:5.886 1st Qu.: 45.02 1st Qu.: 2.100
## Median :0.5380 Median :6.208 Median : 77.50 Median : 3.207
## Mean   :0.5547 Mean   :6.285 Mean   : 68.57 Mean   : 3.795
## 3rd Qu.:0.6240 3rd Qu.:6.623 3rd Qu.: 94.08 3rd Qu.: 5.188
## Max.   :0.8710 Max.   :8.780 Max.   :100.00 Max.   :12.127
##      rad      tax      ptratio      black
## Min.   : 1.000 Min.   :187.0 Min.   :12.60 Min.   : 0.32
## 1st Qu.: 4.000 1st Qu.:279.0 1st Qu.:17.40 1st Qu.:375.38
## Median : 5.000 Median :330.0 Median :19.05 Median :391.44
## Mean   : 9.549 Mean   :408.2 Mean   :18.46 Mean   :356.67
## 3rd Qu.:24.000 3rd Qu.:666.0 3rd Qu.:20.20 3rd Qu.:396.23
## Max.   :24.000 Max.   :711.0 Max.   :22.00 Max.   :396.90
##      lstat      medv
## Min.   : 1.73 Min.   : 5.00
## 1st Qu.: 6.95 1st Qu.:17.02
## Median :11.36 Median :21.20
## Mean   :12.65 Mean   :22.53
## 3rd Qu.:16.95 3rd Qu.:25.00
## Max.   :37.97 Max.   :50.00
```

1. Simple Linear Regression

Fitting the model

```
lm.fit = lm(medv ~ lstat, data = Boston)
```

Calling `summary` we can see information about the fitted model:

- The minimum, maximum and quantile values for the residuals.
- The estimated values for the coefficients, as well as their standard error, and the T-statistic and p-value for the significance test.

- The residual standard error.
- The value of multiple R^2 and adjusted R^2 .
- The value of the model F-statistic and its associated p-value. This statistic measures the relationship between the predictors and the response. When no relationship exists, the F-statistic is expected to be close to 1, whereas it would take values much greater than 1 when this relationship exists.

```
summary(lm.fit)
```

```
##
## Call:
## lm(formula = medv ~ lstat, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.168  -3.990  -1.318   2.034  24.500
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 34.55384    0.56263   61.41  <2e-16 ***
## lstat       -0.95005    0.03873  -24.53  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.216 on 504 degrees of freedom
## Multiple R-squared:  0.5441, Adjusted R-squared:  0.5432
## F-statistic: 601.6 on 1 and 504 DF,  p-value: < 2.2e-16
```

The fitted linear model has the following components:

```
names(lm.fit)
```

```
## [1] "coefficients" "residuals"      "effects"        "rank"
## [5] "fitted.values" "assign"         "qr"            "df.residual"
## [9] "xlevels"      "call"          "terms"         "model"
```

We can print a 95% confidence interval for the coefficients using `confint` (the `level` argument defines the confidence level, defaults to 0.95):

```
confint(lm.fit)
```

```
##              2.5 %      97.5 %
## (Intercept) 33.448457 35.6592247
## lstat       -1.026148 -0.8739505
```

Predictions

When `newdata` is not specified, predictions are done using the fitting (training) data.

```
lm.pred = predict(lm.fit)
```

If we want to pass a list of samples for which we want to predict values:

```
newdata = data.frame(lstat=c(2, 30))
lm.pred2 = predict(lm.fit, newdata = newdata)
lm.pred2
```

```
##      1      2
## 32.65374  6.05236
```

The `interval` argument for `predict` generates intervals for the predicted values. `confidence` returns the 95% *confidence* intervals for the prediction (only reducible error), while `prediction` returns *prediction* intervals considering both reducible and irreducible errors.

```
predict(lm.fit, newdata = newdata, interval = 'confidence')
```

```
##          fit          lwr          upr
## 1 32.65374 31.678068 33.629416
## 2  6.05236  4.625004  7.479716
```

```
predict(lm.fit, newdata = newdata, interval = 'prediction')
```

```
##          fit          lwr          upr
## 1 32.65374 20.402836 44.90465
## 2  6.05236 -6.242765 18.34749
```

Composing Features

Operations over the predictors can be done:

```
lm.fit2 = lm(medv ~ log(lstat), data=Boston)
summary(lm.fit2)
```


```
##
## Call:
## lm(formula = medv ~ log(lstat), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -14.4599  -3.5006  -0.6686   2.1688  26.0129
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  52.1248     0.9652   54.00  <2e-16 ***
## log(lstat)  -12.4810     0.3946  -31.63  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.329 on 504 degrees of freedom
## Multiple R-squared:  0.6649, Adjusted R-squared:  0.6643
## F-statistic: 1000 on 1 and 504 DF, p-value: < 2.2e-16
```

Plotting the data


```
par(mfrow=c(1,2))
plot(lstat, medv, pch='+', cex=.75, title('Linear model'))
abline(lm.fit, col='red', lwd=1.5)

plot(lstat, medv, pch='+', cex=.75, title('Logarithmic model'))
lines(sort(lstat), fitted(lm.fit2)[order(lstat)], col='red')
```

Linear model

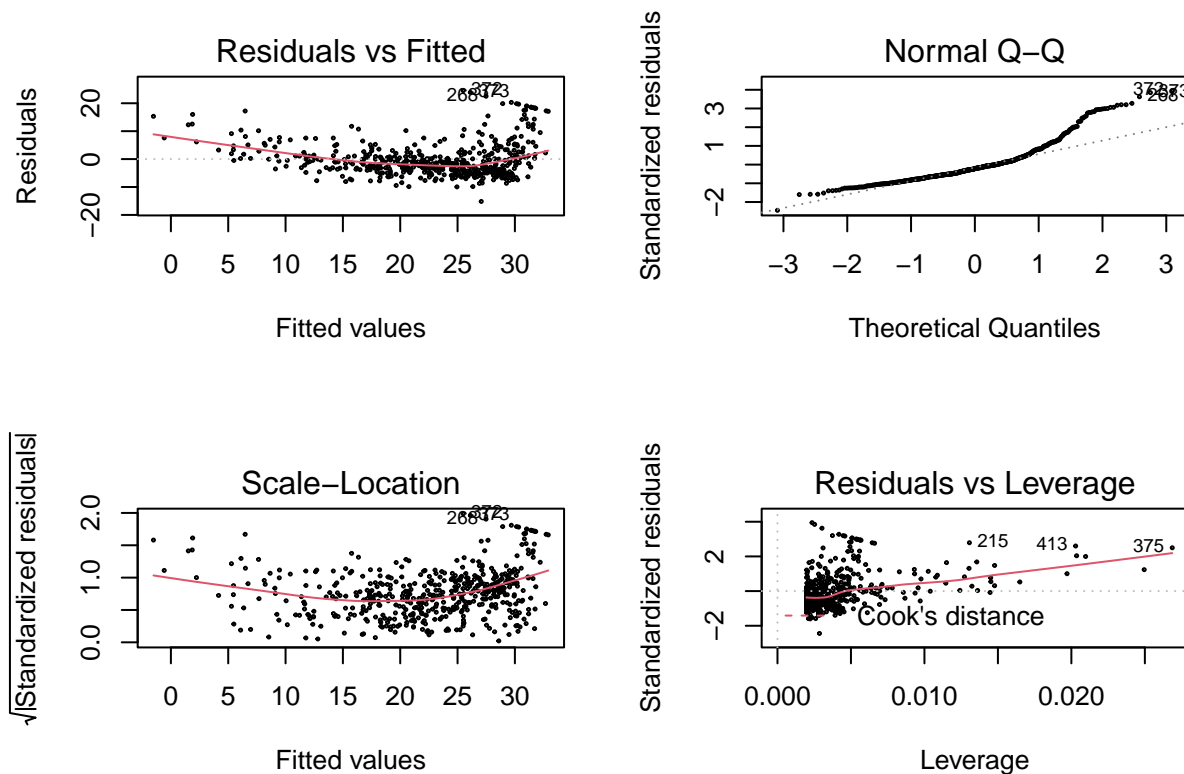


Logarithmic model



The `lm` method comes with a pre-configured 2x2 plot:

```
par(mfrow=c(2,2))  
plot(lm.fit, cex=.25)
```



- The first plot shows Residuals vs Fitted values. This can give an idea of the deviation from linearity by observing the residuals for the predicted values. It's equivalent to `plot(predict(lm.fit), residuals(lm.fit))`
- The second plot is a Normal Q-Q plot, that shows the difference between the model's residuals and a normal distribution, comparing the theoretical quantiles (the quantiles from a standard normal distribution).
- The third plot is obtained by standardizing the residuals from plot number one. Samples with an standardized residual greater than 3 could be considered an outlier. It's equivalent to `plot(predict(lm.fit), rstudent(lm.fit))`
- The fourth plot shows the residual vs the leverage of the sample points. It's equivalent to `plot(hatvalues(lm.fit), rstudent(lm.fit))`.

2. Multiple Linear Regression

Any kind of combination can be applied to features: `poly()`, `log()`, `sqrt()`...

```
lm.mult = lm(medv ~ poly(rm, 2) + sqrt(lstat) + sqrt(rm), data = Boston)
summary(lm.mult)
```

```
##
## Call:
## lm(formula = medv ~ poly(rm, 2) + sqrt(lstat) + sqrt(rm), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -31.0319  -2.6935  -0.4258   2.0009  27.1910
##
```

```
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -341.4959   541.3489  -0.631   0.5284
## poly(rm, 2)1 -415.4564   680.0030  -0.611   0.5415
## poly(rm, 2)2   82.5833    38.1093   2.167   0.0307 *
## sqrt(lstat)   -5.5333     0.2812 -19.675 <2e-16 ***
## sqrt(rm)      152.9910    216.3760   0.707   0.4799
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.609 on 501 degrees of freedom
## Multiple R-squared:  0.7509, Adjusted R-squared:  0.7489
## F-statistic: 377.5 on 4 and 501 DF,  p-value: < 2.2e-16
```

Orthogonal and Non-Orthogonal Polynomials

`poly()` has a `raw` parameter that controls if orthogonal polynomials are created. It defaults to `FALSE`, creating orthogonal polynomials:

```
lm.poly = lm(wage ~ poly(age, 4, raw = FALSE), data = Wage)
coef(summary(lm.poly))
```

```
##              Estimate Std. Error    t value    Pr(>|t|)
## (Intercept)      111.70361   0.7287409  153.283015 0.000000e+00
## poly(age, 4, raw = FALSE)1  447.06785  39.9147851  11.200558 1.484604e-28
## poly(age, 4, raw = FALSE)2 -478.31581  39.9147851 -11.983424 2.355831e-32
## poly(age, 4, raw = FALSE)3  125.52169  39.9147851   3.144742 1.678622e-03
## poly(age, 4, raw = FALSE)4  -77.91118  39.9147851  -1.951938 5.103865e-02
```

`poly()` returns a matrix whose columns are a basis of *orthogonal polynomials*, so each columns is a linear combination of the variables `age`, `age^2`, `age^3` and `age^4`.

If we set `raw=TRUE`:

```
lm.poly = lm(wage ~ poly(age, 4, raw = TRUE), data = Wage)
coef(summary(lm.poly))
```

```
##              Estimate Std. Error    t value    Pr(>|t|)
## (Intercept)      -1.841542e+02  6.004038e+01 -3.067172 0.0021802539
## poly(age, 4, raw = TRUE)1   2.124552e+01  5.886748e+00  3.609042 0.0003123618
## poly(age, 4, raw = TRUE)2  -5.638593e-01  2.061083e-01 -2.735743 0.0062606446
## poly(age, 4, raw = TRUE)3   6.810688e-03  3.065931e-03  2.221409 0.0263977518
## poly(age, 4, raw = TRUE)4  -3.203830e-05  1.641359e-05 -1.951938 0.0510386498
```

Now `poly()` returns `age`, `age^2`, `age^3` and `age^4` directly.

3. Qualitative Predictors

The `Carseats` dataset has both quantitative (numerical) and qualitative predictors:

```
attach(Carseats)
```

Fitting a linear model automatically generates dummy variables for the qualitative predictors:

```
lm.fit3 = lm(Sales ~ ., data=Carseats)
summary(lm.fit3)
```

```
##
## Call:
```

```
## lm(formula = Sales ~ ., data = Carseats)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8692 -0.6908  0.0211  0.6636  3.4115
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   5.6606231   0.6034487   9.380 < 2e-16 ***
## CompPrice     0.0928153   0.0041477  22.378 < 2e-16 ***
## Income        0.0158028   0.0018451   8.565 2.58e-16 ***
## Advertising   0.1230951   0.0111237  11.066 < 2e-16 ***
## Population    0.0002079   0.0003705   0.561  0.575
## Price        -0.0953579   0.0026711 -35.700 < 2e-16 ***
## ShelfLocGood   4.8501827   0.1531100  31.678 < 2e-16 ***
## ShelfLocMedium 1.9567148   0.1261056  15.516 < 2e-16 ***
## Age           -0.0460452   0.0031817 -14.472 < 2e-16 ***
## Education     -0.0211018   0.0197205  -1.070  0.285
## UrbanYes       0.1228864   0.1129761   1.088  0.277
## USYes         -0.1840928   0.1498423  -1.229  0.220
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.019 on 388 degrees of freedom
## Multiple R-squared:  0.8734, Adjusted R-squared:  0.8698
## F-statistic: 243.4 on 11 and 388 DF, p-value: < 2.2e-16
```

In this case, 3 qualitative predictors exist: `ShelveLoc`, `Urban` and `US`.

The model creates one dummy variable for each pair of classes in a predictor. The encoding can be shown using `contrasts()`:

```
contrasts(ShelveLoc)
```

```
##           Good Medium
## Bad           0      0
## Good          1      0
## Medium        0      1
```

For `ShelveLoc`, with 3 classes, 2 dummy variables are created: `ShelveLocGood`, encoded as `[1,0]` and `ShelveLocMedium`, encoded as `[0,1]`; the third class, corresponding to `Bad`, is the trivial encoding, `[0,0]`.